

Architecture Overview:

Here's an overview of what would be on each side of the system... take a look. Lemme know any questions and we can work on how to accomplish this next.

Retailer (satellite) Side:


- **Simple one line embed calls [Widget.js](#) (on our server)** → one-line script embed with Retailer ID calls to a file called [widget.js](#) on our server.. Something like this:

```
<script src="https://mothership.yourdomain.com/widget.js" data-retailer-id="corvette123"></script>
```

- **Captures voice** from customer and sends to mothership server.
- **Displays responses** (voice, text, product info).
- **Retailer SKU Database** remains with retailer.
 - Exposed via **REST API** for mothership queries.

Mothership Side:

- **Central Server** → core application logic.
- **Retailer Preferences DB** → saves branding, tone, voice, UI configs.
- **OpenAI API calls** → Whisper (ASR), GPT (conversation), Voices (TTS).
- **Connector Layer** → queries Retailer SKU DB via REST API.
- **Response** → sends back to retailer widget (voice, text, images, product links).



This is good shit my fren....a recipe....it lines up with what we talked about (satellite widget, mothership brain, SKU connections). i do have a thought again on satellite structure. I'll write it up when I'm a little more rested...But good to have this written down....and yes, I wish it was going faster so we could get to all of it.

Just to give you an idea where I'm at and my new strategy(which might be the way your supposed to do it): I changed from "one thing at a time" to working on multiple things at once when committing qnd deploying. I know that can sound messy, but it feels more efficient when fighting with repos that don't want to behave especially when the right LLM is helping.

Here's my focus right now though(9/19/2025):

- Targeting OpenAI's Realtime Whisper for STT, the same way ChatGPT does it(i think).
- Trying to hook into the exact moment when Whisper detects speech, and then the Realtime TTS kicks in.
- Once that happens, I just want to trigger a simple color change in the button, so the UX can see the shift from idle....showing an active VoxTalk speaking.
- Later, I'd like to replace that with a polished CSS animation (swirl, glow, etc eyc) but after i get a basic fundamental demo completed.

I'm also weighing whether to do something similar when ASR fires (UI talking). Chat doesn't show it that way — UI just stays idle — but if we think it helps, we can hook into the ASR start event too. Maybe it's overkill, maybe not. Either way, the main point is I'm tapping into these events on both the backend and frontend to get a basic demo loop running.

Once that's stable, we can add the “accessories”filters, SKU connectors, etc. but then have more time to do more admin stuff.